

Preliminary report on the emotion classifier project

Dmytro Dumanskyy Tiago Kienzle Hagen Aleksandar Neshov Daniel Pfefferkorn

{dmytro.dumanskyy, t.hagen, a.neshov, d.pfefferkorn}@campus.lmu.de

Abstract

In this report we share our approach to the facial emotion recognition problem. We identified a number of fitting datasets that were well documented in literature and have already been shown to produce good results. Although finding high quality datasets is essential, preprocessing said data is just as important for high quality emotion classification Models. This step is done by comparing all the collected samples and removing duplicates, unusable images where the subject is not clearly recognizable or mislabeled data. Finally the model architecture will be determined after evaluating the accuracy of different, self implemented designs.

1. Introduction

Facial Expression Recognition (FER) is pivotal for Human-Computer Interaction, with critical applications spanning biomedical engineering, healthcare, crime detection, and education [6]. While humans interpret these cues intuitively, replicating this capability computationally remains a significant research challenge.

The primary objective of this project is to implement and evaluate Deep Learning architectures trained from scratch to classify six core emotions—anger, disgust, fear, happiness, sadness, and surprise—using constrained 64×64 pixel inputs. Our initial research found out the necessity of a systematic quality filtering strategy, incorporating techniques such as Laplacian blur detection [11], to effectively process noisy benchmarks like RAF-DB. This report documents our methodology, architectural selection, the integration of Explainable AI (XAI) for the final demonstration, and the key insights gained during this preliminary phase.

2. Datasets

To ensure reliable and comparable results, four widely used benchmark datasets from the literature were selected to train and evaluate the proposed models. All datasets contain facial images categorized into the six target expressions: anger, disgust, fear, happiness, sadness and surprise. The datasets employed are Facial Expression Recognition

2013 Plus (FER+) [2], Real-World Affective Face Database (RAF-DB) [9], Cohn–Kanade (CK+) [10], and AffectNet [12].

FER+ was chosen instead of the original FER2013 due to its improved labelling quality, which has a direct impact on model performance and reliability. In addition, the selected datasets provide diversity in terms of age, ethnicity, and gender, which is essential for improving the generalization capability of the models. Both AffectNet, CK+ and FER+ include an additional contempt label, however, this class is excluded from the experiments, as it falls outside the scope of the present study.

All images have to be resized to 64×64 pixels and grayscaled to facilitate structural exploration and reduce computational complexity. The data includes both grayscale and RGB images. Image formats vary between JPEG and PNG, which does not affect the preprocessing or training pipeline.

Regarding dataset characteristics, the selected subset of AffectNet contains 27,755 RGB images with a resolution of 96×96 pixels in JPEG format. The CK+ dataset consists of 927 grayscale images with a resolution of 48×48 pixels in PNG format. FER+ includes 69,263 grayscale images at 112×112 pixels, also in PNG format. Finally, RAF-DB comprises 15,339 images, both grayscale and RGB, with a resolution of 100×100 pixels in JPEG format.

3. Methodology: Data Preprocessing

We prioritize a large dataset, but try to maintain quality. We clean unconstrained images to remove noise that hurts model training.

3.1. Detection and Geometric Filtering

We use MediaPipe for face detection, increasing the size of low-resolution images using SRGAN [8] for accuracy, because MediaPipe was trained on bigger images.

Occlusion: Images are removed if no face is found or if key face parts (both eyes, nose) are blocked.

Pose Constraints: We reject images with a Yaw (left/right) rotation over 45 degrees. Pitch (up/down) and Roll (tilt) are also limited to avoid extreme angles, ensuring features are visible.

3.2. Backgrounds and Image Quality

Cropping: If the face is small ($< 10\%$ of the image), we crop to center it. We keep some background to help the model handle environmental noise.

Channel Standardization: We convert any single-channel grayscale images to 3-channel mode to ensure dimensional consistency with the model input, though the visual appearance remains unchanged.

Blur & Deduplication: We remove blurry images using the **Laplacian method** [1], this will also remove cartoon images and we remove exact duplicates using **SHA-256 hashing**, or **Perceptual Hashing** [5] if more performance is needed.

3.3. Imbalance and Reduction of Background Noise

To reduce the effect of background noise, we will experiment using an **Attention-Hybrid CNN**. This architecture learns to focus on facial features over background noise. We handle the imbalance in emotion classes using **Class-Balanced Focal Loss** [3] instead of reducing the data, further more about it.

4. Model Architectures

To address the facial emotion recognition (FER) task, our group investigates multiple convolutional neural networks (CNNs) that have demonstrated strong performance in image classification and facial analysis tasks. CNNs build the base for most of our models. This decision is based on according literature emphasizing the clear performance advantage of CNNs compared to other approaches such as traditional machine learning models [4]. Each team member focuses on implementing and evaluating selected architectures in order to identify the most suitable model for achieving a high classification accuracy while maintaining robust generalization performance.

Maximizing the accuracy of the FER task represents the core goal of our group. Apart from the aforementioned aspects, the actual architecture of our neural network plays a key role in achieving this goal to the best of our ability. In the following section we want to highlight our decision making process.

4.1. Criteria for model selection

The suitability of a model for this project is evaluated based on the following criteria:

Firstly the design should be well documented in the literature to ensure scientifically proven effectiveness. For a model to be taken into closer consideration it should have reported strong results on standard benchmarks.

Secondly the architecture should be able to learn complex facial contexts that contribute to the classification of

each emotion. The model has to efficiently recognize important parts of the image such as the eyes, the mouth and the eyebrows and correctly prioritize them over different features of the input such as the background.

Finally models that are prone to overfitting are seen as unfavorable for our project. Given the smaller data quantity relative to other classification tasks such as ImageNet the chosen architecture should be known to generalize well even when trained on a smaller number of images similar to the amount of training samples at our disposal.




Figure 1. This figure from Kopalidis et al. [6] was used as a general reference as to which models to take into consideration.

We took inspiration from the figure above and are eager to try implementing some of the promising networks shown above. The best fitting models for our tasks are the ones that report the highest accuracy while needing proportionally few operations. The Amount of parameters has not been used as a decision metric since especially more modern designs reach high accuracies while not increasing the parameter count drastically.

4.2. Excluded model architectures

Not all types of models were categorized as suitable for the task at hand. Vision Transformers (ViT) are highly data hungry and typically rely on large-scale pretraining, making them suboptimal for training from scratch on limited FER datasets. Furthermore any outdated neural networks that are no longer competitive in terms of accuracy or efficiency were used as baselines for comparison.

4.3. Training setup

We decided to research common training parameters used in scientific papers that yielded good results. All models are trained from scratch using a consistent training protocol to ensure fair comparison.

We decided to use cross entropy loss as our loss function since it's widely used in similar classification tasks. Using stochastic gradient descent (SGD) as our optimizer we are going to train the networks for approximately 30-50 epochs depending on individual model convergence. Furthermore we plan on exploring attention mechanisms to improve robustness against facial occlusions.

To ensure a consistent baseline for model evaluation all models are trained and tested on a predefined dataset.

4.4. Model selection strategy

We are currently in the process of deciding on a final neural network design for the entire group. Each team member is implementing a number of different models based on the architectures that have been proven to achieve good results in FER tasks such as ResNet, EfficientNet, GoogLeNet, Poster++ and others [4], [14]. The final model will be decided through the highest evaluation accuracy on our collective evaluation set.

5. Explainable AI

In order to portray the model's regions of interest in the image we plan on using GradCAM [13], which is a common method used for visualizing and interpreting networks by highlighting the image regions that contribute most strongly to a model's prediction.

6. Demo Implementation

To demonstrate the practical application of our work, we will implement a real-time video pipeline using **OpenCV**, adapting the approach described by Kumar [7].

The planned system consists of two main steps:

1. **Face Detection:** We will use OpenCV's Haar cascade classifiers to scan the video frame. This algorithm identifies the coordinates of the face, allowing us to isolate the facial region from the background.
2. **Classification:** The detected face region will be cropped and resized to 64×64 pixels. This processed image is then passed to our model to classify the emotion.

The final output will display the video feed with a bounding box and the emotion label overlaid in real-time.

References

- [1] Raghav Bansal, Gaurav Raj, and Tanupriya Choudhury. Blur image detection using laplacian operator and open-cv. In *2016 International Conference System Modeling & Advancement in Research Trends (SMART)*, pages 63–67, 2016.
- [2] Emad Barsoum, Cha Zhang, Cristian Canton Ferrer, and Zhengyou Zhang. Training deep networks for facial expression recognition with crowd-sourced label distribution. In *Proceedings of the 18th ACM International Conference on Multimodal Interaction*, page 279–283, New York, NY, USA, 2016. Association for Computing Machinery.
- [3] Yin Cui, Menglin Jia, Tsung-Yi Lin, Yang Song, and Serge Belongie. Class-balanced loss based on effective number of samples. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 9268–9277, 2019.
- [4] Zhixin Dong, Chengdong Wu, Xiangyue Zhang, Xiaosheng Yu, and Junlin Li. A comparison of cnn-based 2d static facial emotion recognition techniques. *Journal of Physics: Conference Series*, 2832(1):012016, 2024.
- [5] Deepak Kandepet. Detecting similar and identical images using perceptual hashes. Hacker Labs (Mirrored via GitHub Gist by augmt), 2017. Accessed: 2025-12-17.
- [6] Thomas Kopalidis, Vassilios Solachidis, Nicholas Vretos, and Petros Daras. Advances in facial expression recognition: A survey of methods, benchmarks, models, and datasets. *Information*, 15(3), 2024.
- [7] Abhijeet Kumar. Demonstration of facial emotion recognition on real time video using CNN: Python & Keras. *Machine Learning in Action*, 2018. Accessed: 2025-12-17.
- [8] Christian Ledig, Lucas Theis, Ferenc Huszár, Jose Caballero, Andrew Cunningham, Alejandro Acosta, Andrew Aitken, Alykhan Tejani, Johannes Totz, Zehan Wang, and Wenzhe Shi. Photo-realistic single image super-resolution using a generative adversarial network. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, pages 4681–4690, Honolulu, HI, USA, 2017. IEEE.
- [9] Shan Li, Weihong Deng, and JunPing Du. Reliable crowdsourcing and deep locality-preserving learning for expression recognition in the wild. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2017.
- [10] Patrick Lucey, Jeffrey F. Cohn, Takeo Kanade, Jason Saragih, Zara Ambadar, and Iain Matthews. The extended cohn-kanade dataset (ck+): A complete dataset for action unit and emotion-specified expression. In *2010 IEEE Computer Society Conference on Computer Vision and Pattern Recognition - Workshops*, pages 94–101, 2010.
- [11] Atharva Malode. How to check for blurry images in your dataset using the laplacian method. *GeeksforGeeks*, 2025. Accessed: 2025-12-17.
- [12] Ali Mollahosseini, Behzad Hasani, and Mohammad H. Mahoor. Affectnet: A database for facial expression, valence, and arousal computing in the wild. *IEEE Transactions on Affective Computing*, 10(1):18–31, 2019.
- [13] Ramprasaath R. Selvaraju, Abhishek Das, Ramakrishna Vedantam, Michael Cogswell, Devi Parikh, and Dhruv Batra. Grad-cam: Why did you say that? visual explanations from deep networks via gradient-based localization. *CoRR*, abs/1610.02391, 2016.
- [14] Christian Szegedy, Wei Liu, Yangqing Jia, Pierre Sermanet, Scott E. Reed, Dragomir Anguelov, Dumitru Erhan, Vincent Vanhoucke, and Andrew Rabinovich. Going deeper with convolutions. *CoRR*, abs/1409.4842, 2014.