

Final report on the emotion classifier project

Dmytro Dumanskyy Tiago Kienzle Hagen Aleksandar Neshov Daniel Pfefferkorn
{dmytro.dumanskyy, t.hagen, a.neshov, d.pfefferkorn}@campus.lmu.de

Abstract

Hello World

1. Introduction

How do humans communicate?

Most people would immediately think of spoken or written language. Yet communication goes beyond words. An essential component of human interaction is *emotion*. Emotions form a key part of nonverbal communication and have been widely studied within the field of psychology.

1.1. What is facial emotion recognition

Facial emotion recognition (FER) refers to the ability to identify human emotions computationally based on images. Other types of emotion recognition tasks approach the same problem by analyzing speech or writing to reconstruct a person's emotional state. As past research has shown, understanding human emotions is no trivial task for an algorithm. The best results in the field of FER have been achieved using complex model architectures and significant amounts of training data. However, the results are quite promising. Recent publications have exceeded accuracies of 90% [Citation] in their studies.

1.2. Why is FER important

Given that humans themselves are capable of emotion recognition, it may not be immediately clear why research in this field is so important. The answer to this question is that characterizing a person's feelings is a key part of many important sectors such as:

- Healthcare
- Public safety
- Crime detection
- Advertising

[Citation]

By automating the emotion recognition process needed in these areas, the risk of human error can be significantly reduced. Furthermore individuals differ substantially in their ability to correctly classify sentiments. Implementing

automated emotion recognition solutions therefore provides a consistent baseline for real-world applications.

1.3. Our work

The goal of this project was to research and understand this technology, followed by the construction, training and evaluation of our own AI model specialized in assessing six basic human emotions: happiness, anger, sadness, fear, disgust and surprise.

The following paper will illustrate the approaches we took to optimize our implementation for accuracy.

2. Related Work

2.1. Efficient Architectures and Low-Resolution Vision

The evolution of computer vision architectures has largely been driven by the quest for efficiency and scalability. Tan and Le introduced **EfficientNetV2** [25], highlighting the importance of training speed and parameter efficiency through the use of Fused-MBConv layers. While highly effective on ImageNet, standard CNNs often struggle with information loss when applied to small inputs. Addressing this, Hassani et al. proposed **Compact Convolutional Transformers (CCT)** [12], demonstrating that transformers can escape the “big data paradigm” by using convolutional tokenizers to handle small datasets effectively. Similarly, modern pure ConvNets like **ConvNeXt V2** [29] utilize fully convolutional masked autoencoders to improve representation learning. More recently, **Vision Mamba (Vim)** [31] challenged the transformer dominance by utilizing bidirectional State Space Models (SSMs) for linear complexity scaling, offering a computationally efficient alternative for visual representation. Our work evaluates these state-of-the-art architectures, adapting their downsampling strategies to the constraints of 64×64 facial expression recognition (FER).

2.2. Transfer Learning in Facial Analysis

While generic object recognition models provide a strong baseline, FER benefits significantly from domain-specific pre-training. Transfer learning from massive face recognition datasets (e.g., CelebA, MS-Celeb-1M) is a standard

practice to overcome the data scarcity in emotion recognition. Research indicates that models pre-trained on millions of faces learn robust, identity-invariant features that are highly transferable. Wang et al. demonstrated this in their **QCS (Quadruplet Cross Similarity)** network [27], utilizing an InsightFace ResNet-50 (IR50) backbone pre-trained on MS-Celeb-1M to extract fine-grained facial features. Inspired by their finding that pre-trained FR backbones provide superior initialization, we adopt the IR50 weights from their repository as a starting point. However, we address a critical limitation: the massive classification heads designed for thousands of identities lead to rapid overfitting on smaller emotion datasets. We propose replacing these heavy heads with lightweight, attention-based mechanisms to preserve spatial information without parameter bloat.

2.3. Lightweight Models and Transfer Learning Strategies for Efficient FER

Recent advances in facial expression recognition (FER) have largely been driven by increasingly deep and parameter-heavy convolutional neural networks. However, the CERN [10] architecture challenges this trend by demonstrating that state-of-the-art (SOTA) performance can be achieved with only 1.45 million parameters. Despite its compact design, CERN outperformed several larger and more computationally demanding models across benchmark datasets. On in-the-wild datasets RAFDB and FerPlus CERN reports 87.09

2.4. Optimization and Generalization Techniques

Training deep networks on heterogeneous, imbalanced datasets requires specialized optimization strategies to ensure convergence and generalization.

- **Warm-up and Learning Rate Schedules:** As demonstrated by Goyal et al. [11], utilizing a learning rate warm-up period is critical for stabilizing training, particularly when using large minibatches. This technique prevents the model from diverging due to large gradient updates during the initial epochs when weights are random.
- **Sharpness-Aware Minimization (SAM):** Standard optimizers like AdamW focus solely on minimizing the training loss value. Foret et al. introduced **SAM** [9], which simultaneously minimizes loss value and loss sharpness. By seeking parameters that lie in “flat minima,” SAM significantly improves generalization performance on out-of-distribution data, a crucial property for our multi-source dataset that includes both lab-controlled and in-the-wild images.

2.5. Choosing a FACS-Aligned Benchmark Dataset for Robust Generalization Testing

CK+ is a well-established dataset focusing on prototypical and micro-expressions, with annotations grounded in the Facial Action Coding System (FACS) developed by Ekman and Friesen [6]. FACS defines facial movements in terms of Action Units (AUs), which correspond to specific underlying muscle activations. The selection of key facial landmarks and their configuration in CK+ follows the FACS framework [18], making it particularly suitable for evaluating models designed to capture fine-grained facial expression dynamics.

3. Methodology

3.1. Datasets

To ensure reliable and comparable results, four widely used benchmark datasets from the literature were selected at the beginning of this study. All datasets contain facial images categorized into the six target expressions: anger, disgust, fear, happiness, sadness, and surprise. The datasets employed were a smaller version of AffectNet [1] (22,629 images), CK+ [3] (927 images), FER+ [8] (56,269 images), and RAF-DB [20] (12,135 images).

FER+ was chosen instead of the original FER2013 due to its improved labeling quality, which has a direct impact on model performance and reliability. Moreover, the selected datasets provide diversity in terms of age, ethnicity, and gender, which is essential for improving the generalization capability of the models. Although AffectNet, CK+, and FER+ include an additional *contempt* label, this class was excluded from the experiments as it falls outside the scope of the present study. The data includes both grayscale and RGB images. Image formats vary between JPEG and PNG, which does not affect the preprocessing or training pipeline.

After the initial training sessions, the developed models exhibited strong overfitting. To mitigate this issue, additional datasets were sought, and access requests were submitted to multiple institutions worldwide to expand the training data.

At the end of the data collection process, a total of 310,215 facial images were gathered. In addition to the previously mentioned datasets, the final corpus also included images from EmoSet [26], JAFFE [15], MMA-FEDB [19], ExpWild [7] and WSEFEP [30]. All datasets were carefully curated and harmonized to match the six target emotion categories used in this study.

Several challenges were encountered during dataset collection and curation. Some datasets contained incorrectly labeled samples, including statues, cartoon characters, or animals assigned to human emotion categories. In other cases, images were organized per subject rather than per

emotion category, without consistent naming conventions, which prevented automated separation. Furthermore, some small datasets were ultimately discarded due to insufficient sample size or poor labeling quality.

As illustrated in the Figure 1, the aggregated datasets are imbalanced, with a significantly higher number of happiness samples compared to the other emotion categories.

3.2. Cross-Dataset Generalization Strategy

For the final training phase, CK+ and KDEF [16] were deliberately excluded from the training set in order to evaluate the generalization capability of the proposed model on unseen datasets. It is important to emphasize that both datasets are relatively balanced.

KDEF also contains high-quality, well-controlled facial images captured under consistent lighting and pose conditions. Although the images were resized to 64×64 pixels to meet the architectural requirements of this study, the dataset remains valuable for assessing cross-dataset robustness due to its standardized acquisition setup.

By excluding CK+ and KDEF from training, these datasets served as independent test sets, enabling a more rigorous evaluation of cross-dataset generalization performance.

3.3. Data Pre-Processing Methodology

The pre-processing pipeline functions as a sequential execution of image evaluation modules, where images failing specific quality thresholds are programmatically discarded. The evolution of this strategy transitioned from strict curation to selective filtering across two empirical phases.

Phase 1: Aggressive Filtering for a Specialized Use Case The primary hypothesis for Phase 1 was that the model would perform optimally if trained exclusively on perfect, high-quality images that mirrored the target demonstration environment (a user looking directly into a high-quality camera). To enforce this pristine standard, the initial pipeline utilized all available filters sequentially:

- **Lighting Filter:** Analyzed the mean pixel intensity and standard deviation, discarding images that were excessively dark, overly bright, or lacked sufficient contrast.
- **Blur Detection:** Assessed image sharpness using frequency analysis via a Fast Fourier Transform (FFT). Images lacking sufficient high-frequency components were rejected.
- **Roll Correction:** Evaluated the image across four rotational states ($0^\circ, 90^\circ, 180^\circ, 270^\circ$), selecting the orientation that yielded the highest face detection confidence.
- **Yaw Filtering:** RetinaFace detected the face, and 6DRepNet predicted the 3D pose. To guarantee a strict frontal viewpoint, any image with a yaw angle exceeding 45° was discarded.

- **Conditional Cropping:** If the ratio of the face area to the total image area was less than 0.4, indicating a small face with significant background, the image was cropped with proportional padding; otherwise, the original image was retained.

In Phase 1, we retained approximately 91,960 images after aggressive filtering.

Phase 2: Resolving the Augmentation Contradiction

The aggressively filtered dataset from Phase 1 caused the models to overfit rapidly. To combat this, data augmentation techniques were applied (adding artificial noise, altering brightness, introducing artificial blur, and utilizing MixUp).

This sequence exposed a fundamental contradiction: we were actively removing naturally difficult images during the pre-processing stage, only to artificially recreate those exact conditions during the augmentation stage. Furthermore, augmentations like MixUp produce overlapping visual structures that standard face detectors completely fail to recognize, meaning such images would be universally rejected by strict filters. Recognizing that the model required this data variance to generalize, Phase 2 drastically reduced the aggression of the filtering pipeline:

- **Reduced Aggression:** The universal application of blur, yaw, roll, and lighting filters was halted. For high-quality datasets (AffectNet, CKplus, FERPlus, RAF-DB), the pipeline was stripped down to utilize only structural standardization.
- **Selective Cleaning:** Aggressive techniques were retained only for uncleaned, chaotic datasets (e.g., ExpW and EmoSet), which utilized the cropping module, face detection validation, and manual exclusion of specific emotion labels.

In Phase 2, we started with 310,215 images. After selective pre-processing, we retained 229,478 images.

Standardization and Deduplication To finalize the datasets, all images underwent structural standardization:

- **Channel Normalization:** All inputs were converted to a 3-channel RGB format to ensure consistent tensor depth.
- **Resizing:** Images were zero-padded to achieve a square aspect ratio and subsequently resized to exactly 64×64 pixels. We utilized cubic interpolation if the original image was smaller than 64×64 .
- **Deduplication:** Duplicate or highly similar images were purged using a combination of perceptual hashing (pHash) and difference hashing (dHash) to ensure strict separation between training and validation distributions.

Dataset Splitting At the beginning, we followed an 80%/10%/10% split strategy, where 80% of the data was used for training, 10% for validation, and 10% for testing. Towards the end of the project, some of us adopted a 90%/10% split strategy in most experiments. The training set comprised 90% of the data, while 10% was reserved for




Figure 1. Number of Images per Emotion After Processing

validation (evaluation). For the final test set, we exclusively utilized two specific high-quality datasets: CK+ and KDEF. This rigorous separation ensured that our final performance metrics reflected true generalization on standardized benchmarks.

3.4. Architectural Adaptations

3.4.1 Basic CNNs and ResNet18 models

To ensure experimental comparability, two baseline convolutional neural networks (3-layer and 5-layer CNNs) were implemented alongside a modified version of ResNet-18, in which the initial 7×7 convolution was replaced with a 3×3 kernel and the input resolution was reduced from 224×224 to 64×64 .

3.4.2 ResNet18 with Squeeze-and-Excitation

Our model is based on ResNet18 and is extended with a squeeze-and-excitation (SE) attention module. ResNet18 serves as a stable baseline with moderate capacity (roughly 11–12 million parameters). Results from compact FER literature, including CERN, suggest that attention mechanisms can improve performance across different model sizes, supporting our use of SE as a low-overhead architectural enhancement [10].

The input to our network is a standardized RGB face image of size 64×64 . Since this resolution is much smaller than the typical ImageNet setup (224×224), we avoid aggressive downsampling in the first layers. In the original ImageNet ResNet, the stem commonly uses a 7×7 convolution with stride 2 followed by max pooling with stride 2. On

a 64×64 image, this reduces spatial resolution very early: 64×64 becomes 32×32 after the first stride-2 convolution, and then 16×16 after max pooling, before the main residual stages even begin. This early reduction removes fine-grained facial details (for example small changes around eyes and mouth) that can be important for emotion recognition. The relationship between kernel size, stride, padding, and output resolution follows the standard convolution output sizing rule [23].

For this reason, we use a 3×3 convolution with stride 1 in the first layer and omit early pooling, so the model preserves more spatial information in the early feature maps. This is also a common modification when ResNet-style models are adapted to small images (e.g., CIFAR-sized inputs), because preserving spatial detail early tends to improve learning on low-resolution inputs.

After the stem, we follow the standard ResNet18 design with four stages of residual blocks and stage-wise downsampling. Residual connections add the block input to the block output, which stabilizes optimization in deep networks [13]. After the final stage, global average pooling produces a feature vector that is fed into a dropout-regularized linear classifier to output logits for the six emotion classes.

We augment the ResNet18 backbone with squeeze-and-excitation (SE) blocks, which perform lightweight channel-wise attention by learning per-channel scaling factors from globally pooled features, improving representational capacity with minimal overhead [14]. The motivation for using a lightweight attention mechanism is consistent with the goal of improving robustness without greatly increasing model size and is aligned with compact FER approaches such as CERN [10].

3.4.3 Loss Function and Class Imbalance

Class imbalance is handled through the loss used for training, which determines how strongly each sample influences the parameter updates. In the standard setting, this loss is cross-entropy. On an imbalanced dataset, cross-entropy can become biased toward majority classes, since they occur more often and therefore contribute a larger share of the gradient updates.

To address this, we modify the training objective by using class-balanced focal loss. The class-balanced component assigns larger weights to underrepresented classes using the effective number of samples [4], while the focal component reduces the contribution of well-classified (easy) samples so that the model continues learning from harder examples [17]. Therefore, we use standard shuffling and address imbalance through the loss rather than through a `WeightedRandomSampler`.

3.4.4 Emotion Prediction Bias and Advanced Loss Functions

A significant challenge encountered in our specific training pipeline was the severe class imbalance within our aggregated dataset. Specifically, images labeled as *happiness* constituted approximately 40% of the total data, while other categories, such as *fear* and *disgust*, were severely under-represented. We sought to implement robust methods to counteract this inherent bias, which otherwise introduces a fundamental model flaw where the majority class becomes the default, high-confidence prediction, while subtle or rare emotions are ignored.

This phenomenon is driven by two factors:

- **Frequency versus Difficulty:** The *happiness* class typically presents distinct visual markers (e.g., a smile) that are easily learned. Conversely, *fear* and *disgust* rely on subtle ocular and brow micro-expressions that degrade heavily at lower resolutions, such as 64×64 pixels.
- **Loss Gradients:** Under standard cross-entropy loss, the model minimizes the global objective by maximizing confidence in the frequent, easy classes, functionally neglecting the minority classes.

To combat this bias, we evaluated two advanced loss functions designed for imbalanced image classification. First, we utilized Class-Balanced Focal Loss (CB-FL), which assigns larger weights to underrepresented classes using the effective number of samples and scales the loss based on prediction confidence, forcing the model to focus on hard, misclassified examples [5]. For EfficientNetV2-S, both advanced loss functions were evaluated. For ResNet-18 SE, only Class-Balanced Focal Loss was applied.

Second, we implemented Label-Distribution-Aware Margin Loss (LDAM) coupled with Deferred Re-Weighting (DRW) [2]. LDAM modifies the logit outputs by assigning a significantly larger mathematical margin to minority classes. This strictly penalizes the model unless it maintains extreme confidence when predicting rare classes, thereby enforcing wider decision boundaries around them. The DRW strategy schedules the training such that the model trains without artificial weights for the initial epochs to learn robust feature representations. In the final epochs, strict class frequency weights are applied to fine-tune the decision boundaries for the minority classes without corrupting the initial feature learning phase.

3.4.5 Transfer Learning (IR50)

We wanted to try bigger models for our project. However, we did not have enough data to train a large model from zero. Transfer learning allowed us to test how this method works in our specific case and opened the possibility to use a larger architecture. A positive side effect was that the pre-trained model (InsightFace ResNet-50) had already learned

facial features from a massive dataset of one million images (Ms-Celeb-1M) [27].

The adaptation occurred in three steps:

Version 1: Original Full Fine-Tuning The first test used the entire IR50 architecture, including the original classification head.

- **The Problem:** The original head flattened the spatial grid into a 25,088-dimensional vector and sent it into a massive linear layer. This single layer had 12.8 million random parameters, accounting for almost half of the whole model.
- **Result:** Training all 27 million parameters on small datasets caused rapid overfitting. The massive linear layer memorized exact spatial positions, and the high learning rate ruined the pre-trained weights.

Version 2: Layer Freezing and the Memorization Problem To protect the pre-trained feature extractors, we froze the early layers (`input_layer`, `body1`, `body2`, and `body3`), which contain about 9 million parameters.

- **Methodology:** We could not freeze the entire backbone because emotion classification requires the network to identify specific facial features that a face recognition model might ignore. Therefore, we kept the final convolutional block (`body4`) trainable.

Version 3: The Light Head Architecture To stop the memorization while keeping `body4` trainable, we built a new head for the model.

- **Methodology:** We removed the 12.8-million-parameter head and replaced it with a “Light Head” containing only 795,000 parameters.
- **Spatial Attention Integration:** Using a standard Global Average Pooling (GAP) layer fixes the parameter problem but destroys spatial information. Human emotions are local, making spatial context critical. We added a Spatial Attention module (Conv1x1 followed by a Sigmoid activation) before the GAP layer. This module learns to give weights to important spatial areas before the GAP layer reduces the dimensions.

3.4.6 Alternative Architectures

During the evaluation phase, three low-parameter architectures were tested to compare different modern deep learning paradigms. Adapting these models for 64×64 inputs required specific architectural modifications, primarily reducing spatial downsampling.

- **Vim (Vision Mamba, ~7M parameters):** Vim [31] was chosen because it is a novel model challenging standard Transformers by utilizing bidirectional State Space Models (SSMs), promising exceptionally low GPU memory usage. To adapt it for 64×64 images, the patch size was changed from 16×16 to 8×8 to provide a functional sequence length of 64 tokens.

- **CCT-7 (Compact Convolutional Transformer, ~4M parameters):** CCT-7 [12] was selected to test a small Transformer specifically designed to handle small images by using a convolutional tokenizer instead of linear patch embeddings.
- **ConvNeXt V2 Atto (~3.7M parameters):** This model [29] was chosen to evaluate how a state-of-the-art pure Convolutional Neural Network operates in a constrained parameter regime. The original architecture features a stem convolution with a stride of 4. For a 64×64 image, this immediately shrinks the feature map to 16×16 . The stem stride was explicitly changed to 1 to preserve spatial information.

3.4.7 EfficientNetV2-S Adaptations

EfficientNetV2-S [25] was selected for its MBConv blocks and Squeeze-and-Excitation (SE) attention, which effectively amplify emotional markers like micro-expressions while suppressing background noise.

Architectural and Augmentation Adaptations To process 64×64 images effectively, two major pipeline adjustments were executed:

- **Stem Stride Fix:** The original stem stride of 2 was reduced to 1. This modification preserved the full spatial resolution through the initial layers, preventing the immediate loss of fine facial details.
- **Minimized Destructive Augmentation:** Initial training with standard ImageNet settings (Mixup $\alpha = 0.8$) obliterated facial features. We implemented a low-interference pipeline to minimize data destruction while maintaining regularization. This included reducing Mixup to 0.2, limiting CutMix to a maximum of 0.3, and using gentle affine transformations. This ensured that while the model was forced to reconstruct features, the core emotional signal remained intact.

3.5. Explainable AI

AI systems that only expose the input and output of the model to the user are called "black box models". The difficulty this design poses lies in the fact that the developer cannot clearly understand the model's reasoning process. Explainable AI aims to solve this problem by providing a method to visualize the importance of different image regions to the final decision of the model. For our task we experimented with two different methods:

3.5.1 Grad-CAM

The grad-cam method is widely used to explain computer vision models. The algorithm works by analyzing the gradients used by the last convolutional layer and assigning a level of importance for each neuron's output for the final classification of the image [21]. However this approach did




Figure 2. Grad-CAM result of the fourth convolutional layer (left) and the third convolutional layer (right)

not seem to work as well as we hoped for our model. Applying grad-cam to our model's last convolutional layer resulted in a coarse and inaccurate heatmap. To combat this problem we experimented with using the previous convolutional layers as the base layer for grad-cam which yielded strongly improved results as in Figure 2.

The poor performance of the grad-cam method is a result of the heavy downsampling during our model's forward pass. The resolution (8×8) of the final feature maps lacks sufficient spacial information to effectively reconstruct the importance of the low resolution input. This also explains why we achieved better results on a layer using higher resolution feature maps.

3.5.2 Integrated gradients

Integrated gradients is a different algorithm to assess the importance of different image regions for the model's decision. This method works by attributing importance values to the pixels of a baseline image. This image is slowly transformed into the original input image and in each transformation step the gradients of each pixel are calculated. The Importance of each pixel is then set as the average gradient of each pixel, which represents how much the model's output changes based on that piece of information [24]. The result is a fine grained attention map which highlights the most important parts of the image as seen in Figure 3.

This approach seemed interesting to us as the quality of the output did not rely on the architectural details of our model.

4. Experiments and Results

4.1. Basic CNNs and ResNet18 Configuration Enhancements

The performance of ReLU and LeakyReLU activation functions was systematically evaluated, with LeakyReLU demonstrating superior overall accuracy. Optimization strategies included Stochastic Gradient Descent (with standard hyperparameters), Adam (with learning rate set as




Figure 3. Heatmap resulting from the integrated gradients method

0.001), and AdamW were tested and Adam with a learning rate of 0.001 yielded the best performance across both the shallow CNN architectures and the modified ResNet-18. Further improvements were achieved through the integration of Batch Normalization, Dropout regularization, and late-stage global average pooling, which enhanced feature stability, mitigated overfitting, and preserved high-level spatial representations relevant to emotion classification.

4.2. ResNet18+SE Ablation: Optimizer and Loss Function

To measure the effect of optimizer choice and loss function while keeping the architecture fixed, we trained the same ResNet18+SE model under four configurations: (1) AdamW + class-balanced focal loss, (2) AdamW + weighted cross-entropy, (3) SGD + class-balanced focal loss, and (4) SGD + weighted cross-entropy. All runs used the same input resolution (64×64), the same augmentation setup, and the same training duration. The validation performance curves are shown in Fig.

Across the four settings, SGD with class-balanced focal loss achieved the strongest validation performance. One likely reason is that AdamW can converge quickly early on, but with noisy and heterogeneous FER data it can also lead to solutions that generalize less well. In addition, class-balanced focal loss tends to be more robust for strong imbalance than weighted cross-entropy because it continues to emphasize hard and minority-class examples later in training [4, 17]. In the final pipeline, imbalance is therefore handled mainly through the loss formulation rather than through sampling-based balancing.

4.3. Transfer Learning (IR50)

Version 1: Original Full Fine-Tuning The first test used the entire IR50 architecture, including the original classification head.

- **The Problem:** The original head flattened the spatial grid into a 25,088-dimensional vector and sent it into a massive linear layer. This single layer had 12.8 million random parameters, accounting for almost half of the whole model.

- **Result:** Training all 27 million parameters on small datasets caused rapid overfitting. The massive linear layer memorized exact spatial positions, and the high learning rate ruined the pre-trained weights.

Version 2: Layer Freezing and the Memorization Problem

- **Result:** We applied different learning rates, but the network still possessed the oversized 12.8-million-parameter head. This setup led to severe overfitting, shown by a performance gap of over 10 percent between training and validation metrics. The massive parameter count allowed the model to memorize the training data.

Version 3: The Light Head Architecture

- **Result:** With only 5.8 million trainable parameters, a reduced MixUp probability, and an Exponential Moving Average (EMA) decay, this setup successfully stopped the overfitting gap and achieved a peak validation accuracy of 81.38 percent.

4.4. Alternative Architectures Evaluation

4.4.1 Vim (Vision Mamba):

However, Vim is fundamentally designed for high-resolution data, and it proved too difficult to make the architecture focus effectively on the fine details of small facial images.

4.4.2 CCT-7 (Compact Convolutional Transformer):

Even though the architecture was perfectly suited for small images, it simply possessed too few parameters to learn the complex features required for human emotion classification.

4.4.3 ConvNeXt V2 Atto:

Similar to CCT-7, the model lacked the parameter count necessary to build robust representations.

4.4.4 MobileNet V2

4.4.5 Other ResNet implementations

Regarding the high accuracy of the ResNet18 + SE implementation ResNet34 and ResNet50 were also trained and evaluated. However the training process was less efficient since the models started overfitting early on due to the high parameter count. As a result the model evaluation yielded a lower accuracy than the original ResNet18 + SE design.

4.5. EfficientNetV2-S Optimization Dynamics

EfficientNetV2-S ultimately emerged as one of the top two best-performing models across all validation benchmarks. To ensure a rigorous comparison, we empirically optimized and tuned each optimizer to find its stable configuration.

Experiment 1: SGD with Nesterov Momentum

- **Result:** Peaked at 74.25 percent validation accuracy.
- **Performance:** SGD [11] proved highly sensitive to hyperparameters. Slight alterations in the learning rate caused massive performance swings, ranging from a complete failure to learn (stagnating below 40 percent accuracy) to achieving a moderate peak of approximately 75 percent.
- **Why it underperformed:** Even at optimal settings, SGD struggled with structural incompatibilities. It applies a uniform learning rate across the network, causing it to under-train depthwise layers while over-training pointwise layers. It lacks the variance tracking required to stabilize the contradictory signals arising from 10 distinct data sources, and it struggles to converge when targets are ambiguous soft labels from Mixup.

Experiment 2: AdamW

- **Result:** Peaked at 83.10 percent validation accuracy.
- **Why it succeeded:** AdamW’s adaptive per-parameter learning rates successfully navigated the heterogeneous architecture, balancing updates between depthwise and dense layers. It provided consistent convergence but exhibited mild overfitting in the final epochs.

Experiment 3: SAM + AdamW (Sharpness-Aware Minimization)

- **Result:** Projected to reach 84–86 percent accuracy.
- **Why it succeeded:** SAM [9] improves generalization by minimizing the worst-case loss in the immediate neighborhood of the parameters. This forces the model into flat, stable minima, which is critical for handling the distributional shifts between our 10 diverse datasets. This robustness allowed it to outperform standard AdamW on the most challenging out-of-distribution test sets (e.g., AffectNet).

Optimizer	Peak Validation Accuracy
SGD	74.25%
AdamW	83.10%
SAM + AdamW	Projected 84–86%

Table 1. Comparison of optimizer performance on EfficientNetV2-S.

4.6. Evaluation of Class Imbalance and Prediction Bias

To quantify the impact of dataset imbalance, we evaluated the EfficientNetV2-S model trained using standard sampling techniques. The validation results exposed a severe disparity in per-class accuracy, confirming the theoretical bias toward visually distinct majority classes.

As presented in Table 2, the model achieved a global accuracy of 60.17%. However, the per-class breakdown reveals that the network massively over-predicted the *happi-*

ness class (96.62% accuracy), while almost entirely failing to recognize *fear* (17.20% accuracy) and *disgust* (29.99% accuracy).

Emotion Class	Accuracy (%)	Correct / Total
Happiness	96.62%	743 / 769
Surprise	83.80%	652 / 778
Sadness	66.71%	483 / 724
Anger	63.36%	472 / 745
Disgust	29.99%	224 / 747
Fear	17.20%	124 / 721
Global	60.17%	2698 / 4484

Table 2. Validation accuracy demonstrating severe prediction bias toward the majority class.

The deployment of LDAM-DRW directly addresses this disparity. By enforcing wider margins for *fear* and *disgust* and deferring the re-weighting until the final training phase, the model is penalized for default *happiness* predictions, yielding a more uniform performance distribution across all six target categories.

4.7. GAN

Inspired by the work with Super-Resolution Generative Adversarial Network [22] we decided to test with a 2x upscaling factor applied at our 64x64 images and then resize it back to 64x64.

Furthermore, a Super Resolution Generative Adversarial Network (SRGAN) method is a GAN that can convert low-resolution images into high-resolution and realistic images. [28]

Unfortunately, due to the image size and that the enhancement process makes the images a bit softer with less facial marks, the training with GAN-enhanced images was worse than with the normal images, achieving easily more than 15% difference between train and evaluation accuracy.

The chosen GAN was GFPGAN, a practical algorithm for real-world face restoration. The model with the second version had the best experimental result and was selected to be applied to all the images, so the group could compare between the normal folder and the GAN-enhanced folder.

4.8. Dual-Input Training

Inspired by the paper about CERN[10] the dual-input strategy was implemented. It consists of giving the training process the same image as two inputs and the correct label. The first image is the normal one and the second is augmented. So in order to receive the correct result, the model has to predict both versions of the image correctly. The results with the dual-input were outperformed by the normal training. It also has increased overfitting, a difference bigger

than 10% (86,45% to 76,15%) between train and evaluation accuracy.

4.9. CERN-Based Architecture

The CERN architecture is derived from a LightCNN backbone with depths ranging from 9 to 29 layers in its full configuration. It employs Max-Feature-Map (MFM) activations for competitive feature selection and integrates Efficient Channel Attention (ECA) to perform lightweight channel-wise recalibration. A central component of the framework is its dual-input training strategy, where the network processes both an original image and its augmented counterpart simultaneously. Correct prediction requires label consistency across both inputs, promoting robustness and invariance to perturbations.

Due to GPU memory constraints, it was not feasible to replicate the complete architecture with combined local and global attention mechanisms. Therefore, a reduced variant consisting of seven convolutional layers and a batch size of 16 was implemented to fit within the available computational resources. Even under this simplified configuration, training remained computationally intensive: 10 epochs required more than 24 hours, and long runs were occasionally interrupted, necessitating continuation from saved checkpoints.

At epoch 47, the model achieved 66.70% evaluation accuracy and 67.51% training accuracy. The learning dynamics showed extended plateaus followed by sudden improvements, suggesting that additional training time could potentially yield further performance gains. However, limited computational capacity—including the use of an NVIDIA RTX 2060 GPU—prevented full exploration of longer training schedules or the use of pretrained weights from the original seven-class model.

5. Conclusion

This project successfully developed a comprehensive emotion classification system for facial expressions. We implemented a complete pipeline from raw image data through preprocessing to model training and evaluation, utilizing state-of-the-art deep learning techniques.

5.1. Key Contributions

1. **Robust Preprocessing Pipeline:** Developed comprehensive preprocessing procedures including face detection, alignment, deduplication, and quality filtering that significantly improved dataset quality and model robustness.
2. **Attention-Enhanced Architecture:** Integrated Squeeze-and-Excitation blocks into ResNet-18, demonstrating consistent performance improvements over baseline models.

3. **Class-Balanced Training:** Employed focal loss to address class imbalance, resulting in more balanced performance across emotion categories.
4. **Multi-Dataset Evaluation:** Comprehensive evaluation across multiple emotion datasets demonstrated model generalization capabilities and robustness across domain shifts.
5. **Model Interpretability:** Applied Grad-CAM visualization to validate that the model learns emotion-specific facial features, providing insight into the decision-making process.

5.2. Limitations and Future Work

While our system achieves competitive performance, several limitations and opportunities for future work exist:

5.2.1 Limitations

- Performance degradation on low-quality images and extreme face angles
- Difficulty distinguishing between similar emotions (fear/surprise, anger/disgust)
- Limited robustness to occlusions and non-frontal faces
- Dependency on extensive preprocessing and high-quality face detection

5.2.2 Future Directions

- Exploration of transformer-based architectures (Vision Transformers, Compact Transformers) for potentially improved performance
- Integration of temporal information for video-based emotion recognition
- Development of robust multi-task learning approaches combining emotion, age, gender, and identity recognition
- Investigation of domain adaptation techniques to improve cross-dataset generalization
- Incorporation of action unit detection for finer-grained emotion understanding
- Real-time deployment optimization for mobile and edge devices

5.3. Final Remarks

The facial emotion classification system presented in this report demonstrates the effectiveness of combining careful data preprocessing, attention mechanisms, and appropriate loss functions for emotion recognition tasks. The model's interpretability through Grad-CAM analysis confirms that it learns meaningful emotional representations. Future work should focus on improving robustness to challenging conditions and exploring more advanced architecture designs, particularly with emerging transformer-based models. The

comprehensive multi-dataset evaluation framework established in this project provides a solid foundation for continued research in this important area of computer vision.

References

- [1] AffectNet. Affectnet dataset. <https://www.kaggle.com/datasets/mstjebashazida/affectnet>.
- [2] Kaidi Cao, Colin Wei, Adrien Gaidon, Nikos Arechiga, and Tengyu Ma. Learning imbalanced datasets with label-distribution-aware margin loss. In *Advances in Neural Information Processing Systems*, 2019.
- [3] Cohn-Kanade. Ck+ dataset. <https://www.kaggle.com/datasets/davilsena/ckdataset>, 2010.
- [4] Yin Cui, Menglin Jia, Tsung-Yi Lin, Yang Song, and Serge Belongie. Class-balanced loss based on effective number of samples, 2019.
- [5] Yin Cui, Menglin Jia, Tsung-Yi Lin, Yang Song, and Serge Belongie. Class-balanced loss based on effective number of samples, 2019.
- [6] Paul Ekman and Wallace V. Friesen. *Manual of the Facial Action Coding System (FACS)*. Consulting Psychologists Press, Palo Alto, CA, USA, 1978.
- [7] Expression in the Wild. Expw dataset. <https://www.kaggle.com/datasets/shahzadabbas/expression-in-the-wild-expw-dataset>.
- [8] FERPlus. Ferplus dataset. <https://www.kaggle.com/datasets/arnabkumarroy02/ferplus>.
- [9] Pierre Foret, Ariel Kleiner, Hossein Mobahi, and Behnam Neyshabur. Sharpness-aware minimization for efficiently improving generalization, 2021.
- [10] Darshan Gera, S. Balasubramanian, and Anwesh Jami. Cern: Compact facial expression recognition net. *Pattern Recognition Letters*, 155:9–18, 2022.
- [11] Priya Goyal, Piotr Dollár, Ross Girshick, Pieter Noordhuis, Lukasz Wesolowski, Aapo Kyrola, Andrew Tulloch, Yangqing Jia, and Kaiming He. Accurate, large minibatch sgd: Training imagenet in 1 hour, 2018.
- [12] Ali Hassani, Steven Walton, Nikhil Shah, Abulikemu Abuduweili, Jiachen Li, and Humphrey Shi. Escaping the big data paradigm with compact transformers, 2022.
- [13] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Deep residual learning for image recognition. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2016.
- [14] Jie Hu, Li Shen, and Gang Sun. Squeeze-and-excitation networks. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.
- [15] JAFFE. Japanese female facial expression (jaffe) dataset. <https://zenodo.org/records/14974867>.
- [16] Karolinska Directed Emotional Faces. Kdef dataset. <https://kdef.se/download-2/register>, 2002.
- [17] Tsung-Yi Lin, Priya Goyal, Ross Girshick, Kaiming He, and Piotr Dollár. Focal loss for dense object detection, 2017.
- [18] Patrick Lucey, Jeffrey F. Cohn, Takeo Kanade, Jason Saragih, Zara Ambadar, and Iain Matthews. The extended cohn-kanade dataset (ck+): A complete dataset for action unit and emotion-specified expression. In *2010 IEEE Computer Society Conference on Computer Vision and Pattern Recognition Workshops (CVPRW)*, pages 94–101, 2010.
- [19] Mahmoud Ima. Mma facial expression dataset. <https://www.kaggle.com/datasets/mahmoudima/mma-facial-expression>.
- [20] RAF-DB. Real-world affective faces database. <https://www.kaggle.com/datasets/shuvoalok/raf-db-dataset>.
- [21] Ramprasaath R. Selvaraju, Abhishek Das, Ramakrishna Vedantam, Michael Cogswell, Devi Parikh, and Dhruv Batra. Grad-cam: Why did you say that? visual explanations from deep networks via gradient-based localization. *CoRR*, abs/1610.02391, 2016.
- [22] Ali I. Siam, Naglaa F. Soliman, Abeer D. Algarni, Fathi E. Abd El-Samie, and Ahmed Sedik. Deploying machine learning techniques for human emotion detection. *Computational Intelligence and Neuroscience*, 2022:8032673, 2022.
- [23] Stanford CS231n. Convolutional neural networks for visual recognition (course notes), 2017. Accessed 2026.
- [24] Mukund Sundararajan, Ankur Taly, and Qiqi Yan. Axiomatic attribution for deep networks. *CoRR*, abs/1703.01365, 2017.
- [25] Mingxing Tan and Quoc V. Le. Efficientnetv2: Smaller models and faster training, 2021.
- [26] VCC Tech. Emoset dataset. <https://vcc.tech/EmoSet>.
- [27] Chengpeng Wang, Li Chen, Lili Wang, Zhaofan Li, and Xuebin Lv. Qcs: Feature refining from quadruplet cross similarity for facial expression recognition, 2025.
- [28] Xiang Wei, Ting Feng, Qinghua Huang, Qian Chen, Chao Zuo, and Haigang Ma. Deep learning-powered biomedical photoacoustic imaging. *Neurocomputing*, 573:127207, 2024.
- [29] Sanghyun Woo, Shoubhik Debnath, Ronghang Hu, Xinlei Chen, Zhuang Liu, In So Kweon, and Saining Xie. Convnext v2: Co-designing and scaling convnets with masked autoencoders, 2023.
- [30] WSEFEP. Warsaw set of emotional facial expression pictures. <https://www.emotional-face.org/hopsa-sa>.
- [31] Lianghui Zhu, Bencheng Liao, Qian Zhang, Xinlong Wang, Wenyu Liu, and Xinggang Wang. Vision mamba: Efficient visual representation learning with bidirectional state space model, 2024.